

UniIO Performance Scripts

- [UniIO Performance Scripts](#)
 - [install](#)
 - [client scripts](#)
 - [client/plotfio.sh](#)
 - [server scripts](#)
 - [server/collect_cpu.sh](#)
 - [server/counterana.py](#)
 - [server/init_backend.sh](#)
 - [server/init_cluster.sh](#)
 - [auto perftest scripts](#)
 - [auto/perfauto.py](#)

install

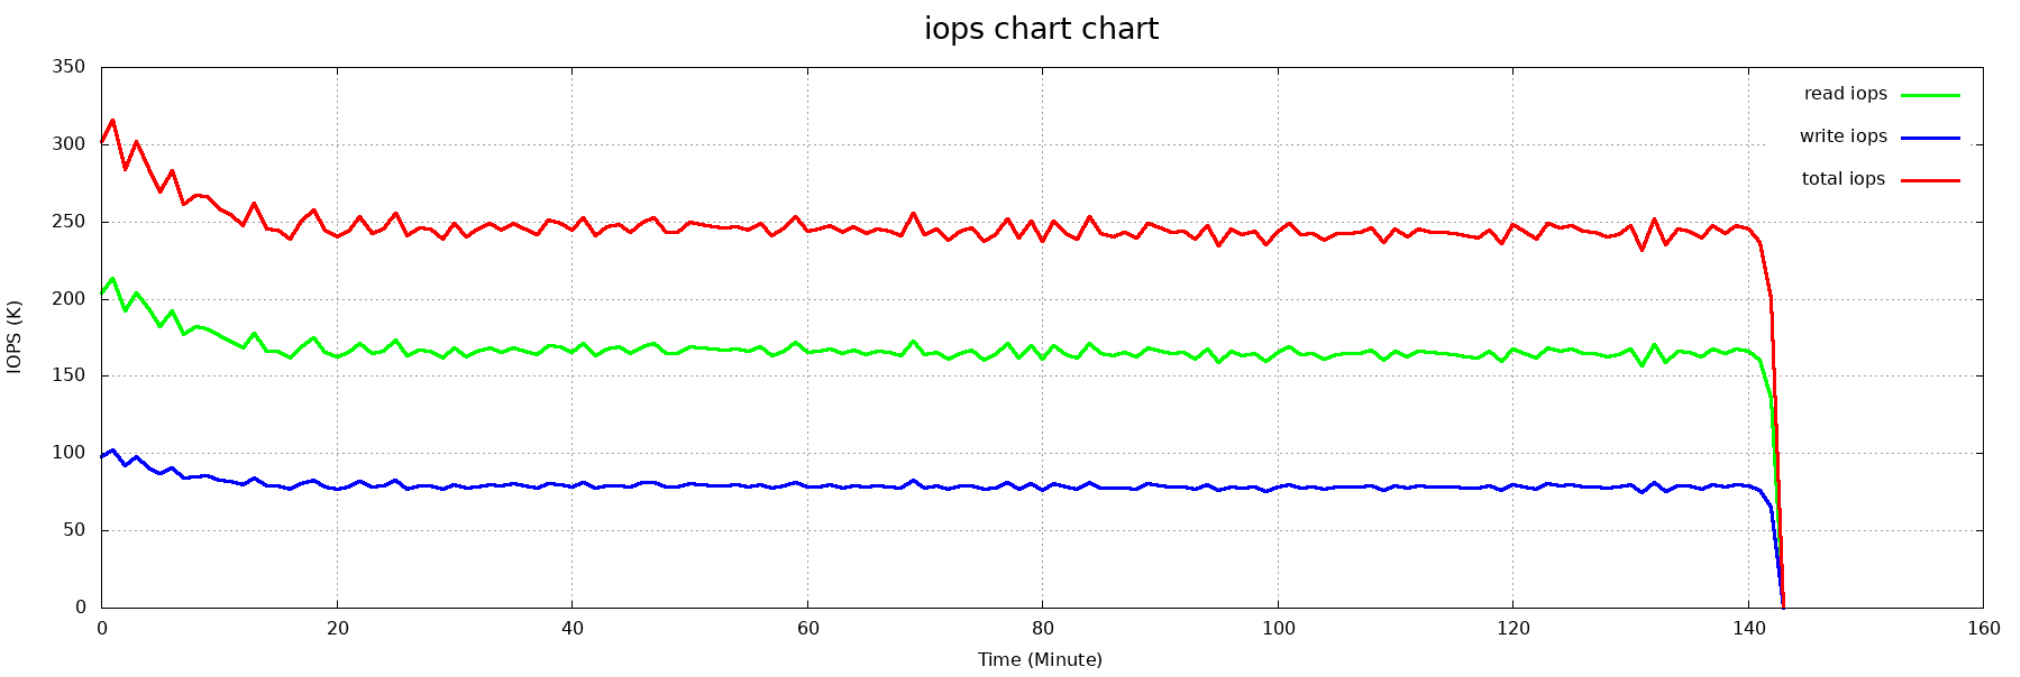
```
$ git clone https://github.com/fred-chen/uio_scripts.git
$ tree uio_scripts/
uio_scripts/
├── client
│   └── plotfio.sh
└── server
    ├── collect_cpu.sh
    ├── counterana.py
    ├── init_backend.sh
    ├── init_cluster.sh
    └── renice_iothreads.sh
```

client scripts

client/plotfio.sh

```
##### fio #####
##### fio ##### job ##### IOPS SLAT, CLAT, LAT
###
# client/plotfio.sh -h
usage: plotfio.sh <logname> [-t iops|clat|slat|lat] [--title chart_title] [-k|--keep]
options:
  -t: type of plots, can be one of: iops, clat, slat, lat.
  -k: keep temp files.
examples:
  plotfio.sh log/82rw*iops* -t iops # plot iops chart for logs that the path match 'log/82rw*iops*'
```

IOPS




server scripts

server/collect_cpu.sh

```
##### linux bcc/eBPF ##### oncpu offcpu wakeup offwakeup ##### SVG #####
#####
```

111



server/counterana.py

```
$ uio_scripts/server/counterana.py -h
```

Analyze UniIO counter log files.

-i: ignore case

sample_count: how many samples(lines) have been aggregated for a counter
unit: the unit of a counter (counts, uSec, KiB)
trends: trends of the sample value from the first sample to the last in [UP|DOWN|FLAT|NOCHANGE|SPIKES]
slop: result of linear regression(the 'a' in $y=ax+b$). how fast the sample value increase|decreases
self explained: min, max, mean, mean_squared_deviation, standard_deviation, pct_stddev:mean

```
.. counter = counter + 1 if direction == 'UP' else counter - 1
```

2. `ss.obs.cacheMigrateFromWriteToRead` (stddev:mean=58.9%) **UP** (-m) counter


3 / 7

```
(2097152...4194304] 1
(4194304...8388608] 2
(8388608...16777216] 3
(16777216...33554432] 8
(33554432...67108864] 15
(67108864...134217728] 30
(134217728...268435456] 58
(268435456...536870912] 108
(536870912...1073741824] 216
(1073741824...2147483648] 81
```

3. 生成counter数据并生成counter数据(-g)数据

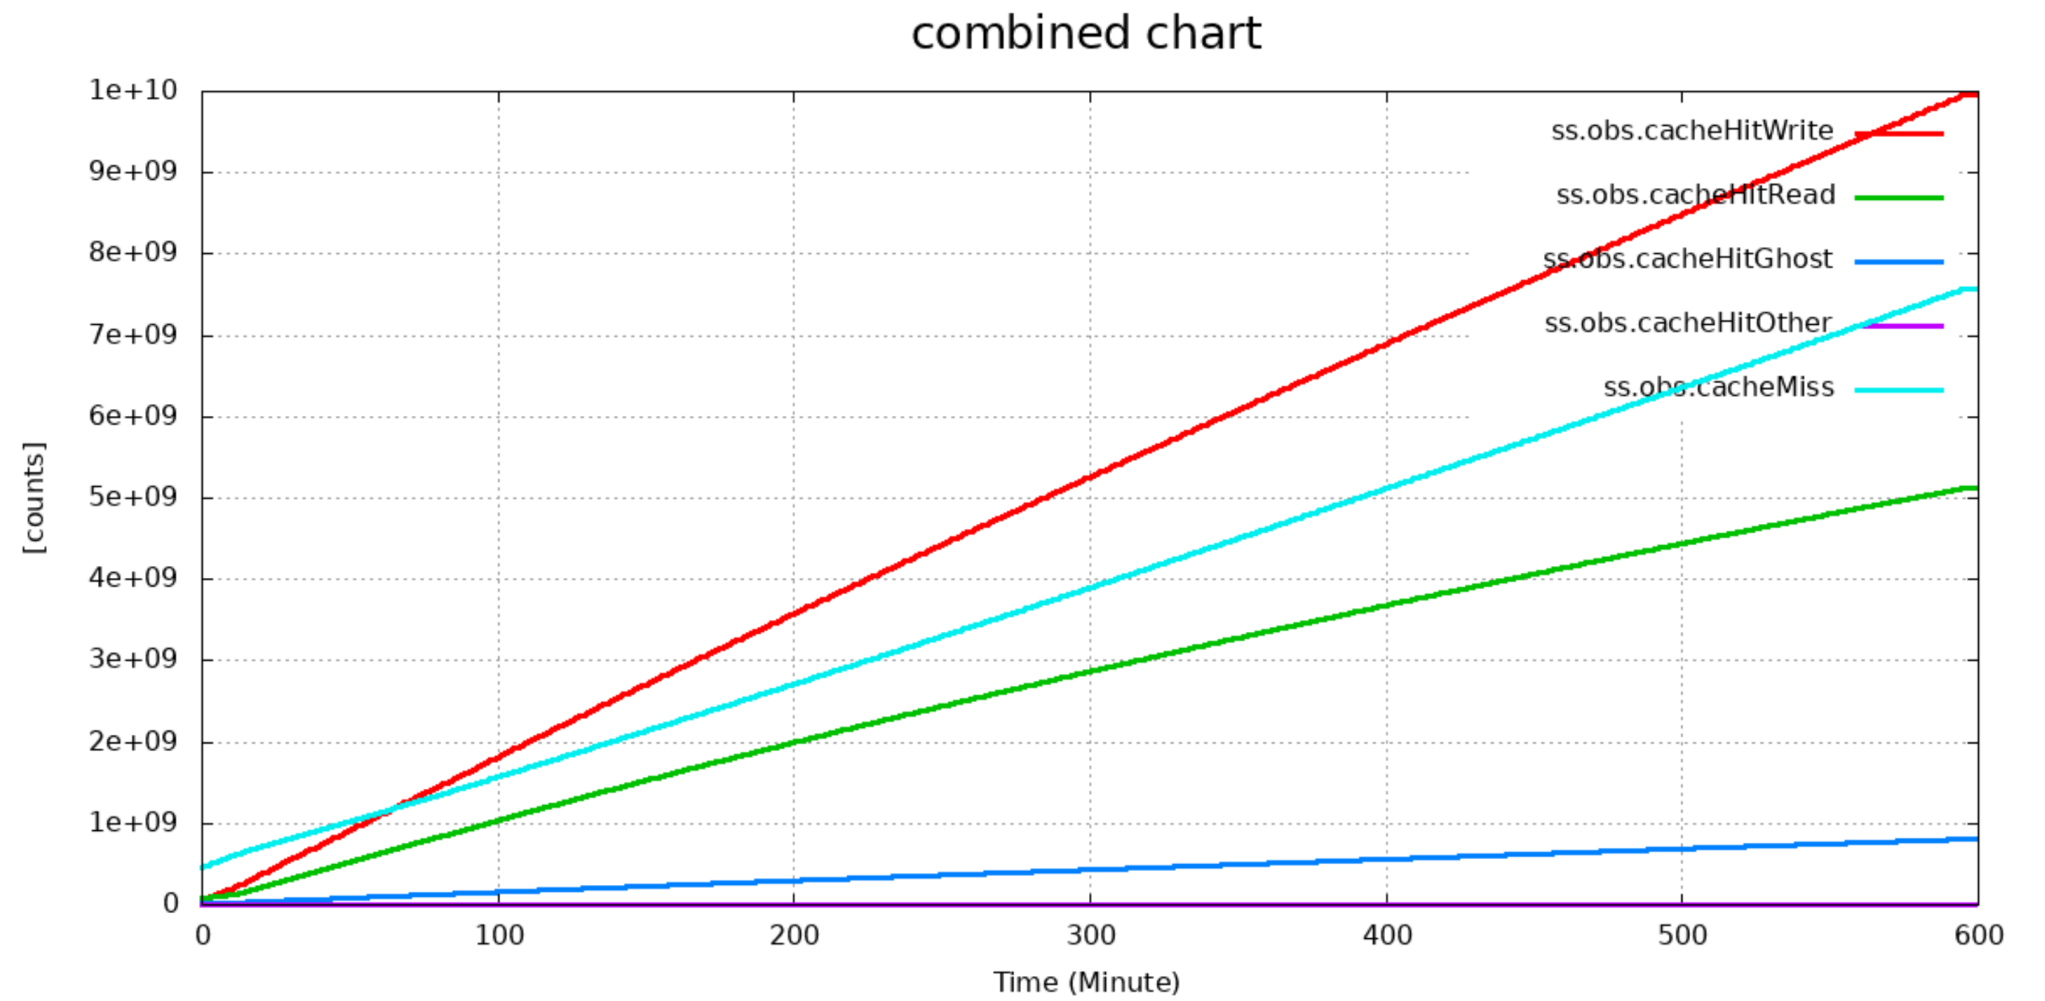
```
$ server/counterana.py counter.log -e ss.obs.cacheMigrateFromWriteToRead -g
building aggregated array ... done.
=====
ss.obs.cacheMigrateFromWriteToRead[523][counts][UP]: min=310342.0 max=1260108523.0 mean=627497530.2 stddev=371270071.4
stddev:mean=59.2% slop=2458905.563
=====
ss.obs.cacheMigrateFromWriteToRead.png
```

4. 生成ss.obs.cacheMigrateFromWriteToRead.png 数据并生成(使用 [plotfio.sh](#) 生成iops/latency)数据



5. 生成counter 数据

```
$ uio_scripts/server/counterana.py counter.log -e 'ss.obs.cache(=?=Miss|Hit)' -gc
building aggregated array ... done.
=====
ss.obs.cacheHitWrite[600][counts][UP]: min=61190410.0 max=9962262664.0 mean=5170040745.2 stddev=2889397492.9
stddev:mean=55.9% slop=16676433.790
ss.obs.cacheHitRead[600][counts][UP]: min=79458824.0 max=5126165414.0 mean=2767213820.0 stddev=1479525548.5
stddev:mean=53.5% slop=8527487.472
ss.obs.cacheHitGhost[600][counts][UP]: min=13269294.0 max=806473330.0 mean=423335945.3 stddev=230776016.9
stddev:mean=54.5% slop=1331940.169
ss.obs.cacheHitOther[600][counts][NOCHANGE]: min=0.0 max=0.0 mean=0.0 stddev=0.0 stddev:mean=0.0% slop=0.000
ss.obs.cacheMiss[600][counts][UP]: min=454818307.0 max=7578260188.0 mean=3942654687.8 stddev=2070806872.7
stddev:mean=52.5% slop=11952484.988
=====
/root/fred/ss.obs.cacheHitWrite_more.plotdata.png
```



UniIO counters

```
$ cat counters.sh
#!/usr/bin/env bash
#usage: ./counters.sh [interval] [runtime]
runtime=36000 # how long, default 10 hours
interval=60 # how often, default every 60 seconds
[[ ! -z "$1" ]] && runtime=$1
[[ ! -z "$2" ]] && interval=$2
total=0
while true
do
    date
    arrayctl counters
    sleep $interval
    total=$((total+$interval))
    [[ $total -ge $runtime ]] && break
done
$ nohup ./counters.sh 36000 > counter.log 2>&1 &
```

server/init_backend.sh

```
1. UniIO
2. DP 'config.ini'
3. coredump
root
$ server/init_backend.sh -h
usage: init_backend.sh [ clear|init ] [ -G dumpdev_size ]
```

server/init_cluster.sh

```
!!!
$ server/init_cluster.sh -h
usage: init_cluster.sh [-f] [-s|--stoponly]
        [-b|--bootonly]
        [-r|--replace rpm_dir]
        [-d|--initbackend] [-G dump_size]
        [-i|--initarray]
        [-c|--createluns --management_ip ip --iscsi_ip ip --topology ip,ip...]
-f: force (killing cio_array)
```

- s: stop only
- b: start objmgr and objmgr-fab
- d: initialize backend
- G: prereserve size for coredump device
- i: initialize array
- c: create new luns and mappings
- management_ip: specify the management IP address for the federation
- iscsi_ip: specify the management IP address for the federation
- topology: specify the node IP addresses for the federation

auto perftest scripts

auto/perfauto.py

```

# UnioIO Federation 脚本, fio 脚本, 性能测试脚本
1. '-c' 创建 UnioIO Federation 脚本 perfauto.py 脚本 uniio, uniio-ui, sysmgmt, nasmgmt RPM 脚本 RPM
2. '-u' 更新 UnioIO Federation 脚本 perfauto.py 脚本 uniio, uniio-ui, sysmgmt, nasmgmt RPM 脚本 RPM
3. '-u --binonly' 更新 RPM 脚本 uniio 脚本 cio_array, cio_array.sym) 更新 UnioIO Federation 脚本
UnioIO Federation 脚本 '/opt/uniio/sbin/cio_array' 脚本 '-u --binonly=./replacefile/cio_array'
4. '--binonly=xxx' 更新 cio_array 脚本 cio_array.sym 脚本 '--binonly=xxx' 更新 git 脚本 commit
git 脚本 commit 脚本 '-c' 更新 git 脚本 '--binonly=xxx' 更新 git 脚本
2. a. uniio RPM 2. 脚本 '--binonly=' 脚本 '--binonly=xxx' 脚本 '--binonly=conf' 脚本 git
5. '-i' 更新 uniio 脚本 'perfauto.py' 脚本 'init_cluster.sh' 脚本
6. '-p' 脚本 fio 脚本 lun 脚本 fio 脚本 counter 脚本 cpu 脚本
'runfio.sh' 脚本 'counters.sh'
7. '-p --cpudata' 脚本 fio 脚本 uniio 脚本 'collect_cpu.sh' 脚本 cpu
8. '-p --fill=sec' 脚本 fio 脚本 LUN 脚本 sec
9. '--fullmap' 脚本 '--createluns' 脚本 ISCSI 脚本 LUN 脚本 LUN
$ uio_scripts/auto/perfauto.py -h
usage: perfauto.py [ -c|--config configfile.json ]
        [ -f|--force ] [ -s|--shutdown ]
        [ -b|--boot ]
        [ -u|--update ] [ --binonly (binpath|conf|tag|branch|commit) ]
        [ -i|--init ]
        [ -p|--perftest ] [ --cpudata ] [ --fill sec ]
        [ --createluns num ] [ --fullmap ] [ --deleteluns ]
Coordinate UnioIO nodes, build server and fio clients for performance test.
options:
-c, --config:      config file path (.json)
-f, --force:       force stop uniio node (kill cio_array)
-s, --shutdown:    gracefully stop uniio nodes
-b, --boot:        start uniio nodes
-u, --update:      update uniio build
    --binonly:     use along with '-u', only update cio_array binary.
-i, --init:        reinit uniio federation
-p, --perftest:    run perftest
    --cpudata:     use along with '-p', collect cpu data as svg files while performance test is running
    --fill:        use along with '-p', fill the luns with pure write workload for a given time in seconds
--createluns:      create a given number of luns
    --fullmap:     use along with '--createluns', all clients see all luns ( clients see different luns if not specified )
--deleteluns:      delete all existing luns

```

配置文件

```

{
  "runtime_dir" : "/tmp/uio",
  "client_nodes" : [
    ["192.168.100.169", "root", "p@ssword"],
    ["192.168.100.155", "root", "password"],
    ["192.168.100.156", "root", "password"]
  ],
  "federation_nodes" : [
    ["192.168.100.206", "root", "password"],
    ["192.168.103.248", "root", "password"],
    ["192.168.101.169", "root", "password"]
  ],
  "build_server" : ["192.168.100.120", "root", ".id_rsa", "/root/fred/.id_rsa"],
  "build_server_git_proxy" : "socks5://192.168.100.120:8899",
  "uniio_checkout" : "default",

```

```
"num_luns" : 18,
"lunsize_G" : 1000,
"topology" : "192.168.101.169,192.168.103.248,192.168.100.206",
"management_ip" : "192.168.103.253",
"iscsi_ip" : "192.168.60.253",
"fio_runtime" : 10800,
"fio_ramp_time" : 0,
"fio_dedupe_percentage" : 80,
"fio_buffer_compress_percentage" : 60,
"fio_random_distribution" : "random",
"### fio_random_distribution can be any fio supported distributions: [random, zipf:0.96, pareto:ratio, ..]" : "",
"fio_rw" : "randrw",
"### fio_rw can be 'sepjob[_fio-supported-rw]' or any fio supported rw types" : "",
"### fio_rw 'sepjob_xxx' means use different jobs for read and write in mixed workload" : "",
"### fio_rw example: 'sepjob', 'sepjob_randrw', 'sepjob_rw' " : "",
"fio_rwmixread" : 80,
"fio_rwmixwrite" : 20,
"runfio_jobs" : "1",
"runfio_qdepth" : "4",
"runfio_xxx is arguments for 'runfio.sh', e.g. --jobs --qdepth" : ""
}
```